


LECTURE 35


Dr.Vani V

Source : Han & Kambar , Chapter 10, Data Mining Concepts &
Techniques, BIRCH & DBScan



BIRCH: Multiphase Hierarchical Clustering Using Clustering Feature Trees

- **Balanced Iterative Reducing and Clustering using Hierarchies** (BIRCH) is designed for clustering a large amount of numeric data by **integrating hierarchical clustering** (at the initial micro-clustering stage) and **other clustering methods such as iterative partitioning** (at the later macro-clustering stage).
- It overcomes the two difficulties in agglomerative clustering methods: (1) scalability and (2) the inability to undo what was done in the previous step.
- BIRCH uses the notions of clustering feature to summarize a cluster, and **clustering feature tree (CF-tree)** to represent cluster hierarchy.




BIRCH: Multiphase Hierarchical Clustering Using Clustering Feature Trees

- The structures help the clustering method achieve good speed and scalability in large or even streaming databases.
- Make it effective for incremental and dynamic clustering of incoming objects.

Consider a cluster of n d -dimensional data objects or points. The clustering feature (CF) of the cluster is a 3-D vector summarizing information about clusters of objects. It is defined as

$$CF = \langle n, LS, SS \rangle,$$

where LS - Linear Sum of the n points
SS – Square Sum of the data points



BIRCH: Multiphase Hierarchical Clustering Using Clustering Feature Trees

- A clustering feature is a summary of the statistics for the given cluster.
- Using a clustering feature, derive many useful statistics of a cluster.
- For example, the cluster's centroid, x_0 , radius, R , and diameter, D , are

$$x_0 = \frac{\sum_{i=1}^n x_i}{n} = \frac{LS}{n},$$

$$R = \sqrt{\frac{\sum_{i=1}^n (x_i - x_0)^2}{n}} = \sqrt{\frac{nSS - 2LS^2 + nLS}{n^2}},$$

$$D = \sqrt{\frac{\sum_{i=1}^n \sum_{j=1}^n (x_i - x_j)^2}{n(n-1)}} = \sqrt{\frac{2nSS - 2LS^2}{n(n-1)}}.$$


- Here, $R \rightarrow$ average distance from member objects to
- $D \rightarrow$ the average pairwise distance within a cluster.
- Both R and D reflect the tightness of the cluster around



BIRCH: Multiphase Hierarchical Clustering Using Clustering Feature Trees

- Summarizing a cluster using the clustering feature can avoid storing the detailed information about individual objects or points.
- Only need a constant size of space to store the clustering feature. This is the key to BIRCH efficiency in space.
- Clustering features are additive. That is, for two disjoint clusters, $C1$ and $C2$, with the clustering features
 - $CF1 = \{n1, LS1, SS1\}$ and $CF2 = \{n2, LS2, SS2\}$, respectively
- The clustering feature for the cluster that formed by merging $C1$ and $C2$ is simply

$$CF_1 + CF_2 = \langle n_1 + n_2, LS_1 + LS_2, SS_1 + SS_2 \rangle$$




BIRCH : Example

- Clustering feature. Suppose there are three points, (2,5), (3,2), and (4,3), in a cluster, C1. The clustering feature of C1 is

$$CF_1 = \langle 3, (2 + 3 + 4, 5 + 2 + 3), (2^2 + 3^2 + 4^2, 5^2 + 2^2 + 3^2) \rangle = \langle 3, (9, 10), (29, 38) \rangle.$$

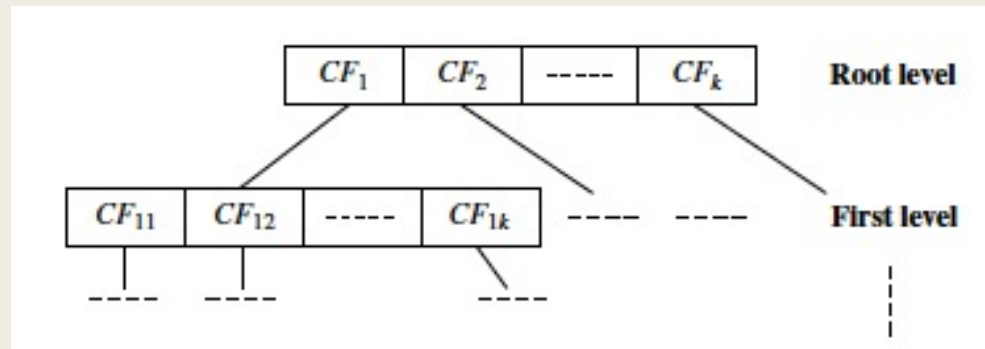
- Suppose that C1 is disjoint to a second cluster, C2, where $CF_2 = \langle 3, (35, 36), (417, 440) \rangle$.
- The clustering feature of a new cluster, C3, that is formed by merging C1 and C2, is derived by adding CF_1 and CF_2 . That is,

$$CF_3 = \langle 3 + 3, (9 + 35, 10 + 36), (29 + 417, 38 + 440) \rangle = \langle 6, (44, 46), (446, 478) \rangle.$$



BIRCH: CF Tree

- A CF-tree is a height-balanced tree that stores the clustering features for a hierarchical clustering.



- By definition, a non-leaf node in a tree has descendants or “children.”
- The non-leaf nodes store sums of the CFs of their children, and thus summarize clustering information about their children.
- A CF-tree has two parameters: branching factor, **B**, and threshold, **T**.
 - The **branching factor** specifies the maximum number of children per non-leaf node.
 - The **threshold parameter** specifies the maximum number of subclusters stored at the leaf nodes of the tree.
 - These two parameters implicitly control the result of the clustering.



BIRCH Phases...

- Given a limited amount of main memory, an important consideration in BIRCH is to minimize the time required for input/output (I/O).
- BIRCH applies a multiphase clustering technique:
 - *A single scan of the data set yields a basic, good clustering, and*
 - *one or more additional scans can optionally be used to further improve the quality. T*
- The primary phases are
 - **Phase 1:** *BIRCH scans the database to build an initial in-memory CF-tree, which can be viewed as a multilevel compression of the data that tries to preserve the data's inherent clustering structure.*
 - **Phase 2:** *BIRCH applies a (selected) clustering algorithm to cluster the leaf nodes of the CF-tree,*




BIRCH Phases...

- For Phase 1, the CF-tree is built dynamically as objects are inserted.
 - *The method is incremental.*
 - *An object is inserted into the closest leaf entry (subcluster).*
 - *If the diameter of the subcluster stored in the leaf node after insertion is larger than the threshold value, then the leaf node and possibly other nodes are split.*
 - *After the insertion of the new object, information about the object is passed toward the root of the tree.*
 - *The size of the CF-tree can be changed by modifying the threshold.*
 - *If the size of the memory that is needed for storing the CF-tree is larger than the size of the main memory, then a value can be specified, and the CF-tree is rebuilt.*




BIRCH Phases

- The rebuild process is performed by building a new tree from the leaf nodes of the old tree. Thus, the process of rebuilding the tree is done without the necessity of rereading all the objects or points.
- This is like the insertion and node split in the construction of **B+-trees**. Therefore, for building the tree, data must be read just once. Some heuristics and methods have been introduced to deal with outliers and improve the quality of CF-trees by additional scans of the data.
- Once the CF-tree is built, any clustering algorithm, such as a typical partitioning algorithm, can be used with the CF-tree in Phase 2.




How effective is BIRCH?

- The time complexity of the algorithm is $O(n)$, where n is the number of objects to be clustered.
- Experiments have shown the linear scalability of the algorithm with respect to the number of objects, and good quality of clustering of the data. However, since each node in a CF-tree can hold only a limited number of entries due to its size, a CF-tree node does not always correspond to what a user may consider a natural cluster.
- **If the clusters are not spherical in shape, BIRCH does not perform well because** it uses the notion of radius c to control the boundary of a cluster.



Density-Based Methods

- Partitioning and hierarchical methods are designed to find spherical-shaped clusters.
- They have difficulty finding clusters of arbitrary shape such as the “S” shape and oval clusters




- Given such data, they would likely inaccurately identify regions, where noise or outliers are included in the



Density-Based Methods


- To find clusters of arbitrary shape,
 - *model clusters as dense regions in the data space, separated by sparse regions.*
 - *This is the main strategy behind **density-based clustering methods**, which can discover clusters of nonspherical shape.*



DBSCAN...

“How can we find dense regions in density-based clustering?”


- *The density of an object o can be measured by the number of objects close to o .*
- *DBSCAN (**Density-Based Spatial Clustering of Applications with Noise**) finds core objects, that is, objects that have dense neighborhoods.*
- *It connects core objects and their neighborhoods to form dense regions as clusters.*



DBSCAN...

“How does DBSCAN quantify the neighborhood of an object?”

- *A user-specified parameter $\epsilon > 0$ is used to specify the radius of a neighborhood we consider for every object.*
- *The ϵ -neighborhood of an object o is the space within a radius ϵ centered at o .*
- *Due to the fixed neighborhood size parameterized by ϵ , the density of a neighborhood can be measured simply by the number of objects in the neighborhood.*
- *To determine whether a neighborhood is dense or not, DBSCAN uses another user-specified parameter, $MinPts$, which specifies the density threshold of dense regions.*
- *An object is a core object if the ϵ -neighborhood of the object contains at least $MinPts$ objects. Core objects*
define dense regions.



DBSCAN...


- Given a set, D , of objects, we can identify all core objects with respect to the given parameters, ϵ and MinPts.
- The clustering task is therein reduced to using core objects and their neighborhoods to form dense regions, where the dense regions are clusters.
- For a core object q and an object p , we say that p is directly density-reachable from q (with respect to ϵ and MinPts) if p is within the ϵ -neighborhood of q .
- An object p is directly density-reachable from another object q if and only if q is a core object and p is in the ϵ -neighborhood of q .
- Using the directly density-reachable relation, a core object can “bring” all objects from its ϵ -neighborhood into a



DBSCAN...


“How can we assemble a large dense region using small dense regions centered by core objects?”

- *In DBSCAN, p is density-reachable from q (with respect to ϵ and MinPts in D) if there is a chain of objects p_1, \dots, p_n , such that $p_1 = q$, $p_n = p$, and p_{i+1} is directly density-reachable from p_i with respect to ϵ and MinPts , for $1 \leq i \leq n$, $p_i \in D$.*
- *Note that density-reachability is not an equivalence relation because it is not symmetric. If both o_1 and o_2 are core objects and o_1 is density-reachable from o_2 , then o_2 is density-reachable from o_1 . However, if o_2 is a core object but o_1 is not, then o_1 may be density-reachable from o_2 , but not vice versa.*





DBSCAN...

- To connect core objects as well as their neighbors in a dense region, **DBSCAN uses the notion of density-connectedness.**
- Two objects $p_1, p_2 \in D$ are density-connected with respect to ϵ and MinPts if there is an object $q \in D$ such that both p_1 and p_2 are density reachable from q with respect to ϵ and MinPts.
- **Unlike density-reachability, density connectedness is an equivalence relation.** It is easy to show that, for objects o_1 , o_2 , and o_3 , if o_1 and o_2 are density-connected, and o_2 and o_3 are density-connected, then so are o_1 and o_3 .





Example: Density-reachability and density-connectivity

Consider Figure for a given ϵ represented by the radius of the circles, and, say, let $\text{MinPts} = 3$.



Example: Density-reachability and density-connectivity


- Of the labeled points, m, p, o, r are core objects because each is in an ϵ -neighborhood containing at least three points.
- Object q is directly density-reachable from m . Object m is directly density-reachable from p and vice versa.
- Object q is (indirectly) density-reachable from p because q is directly density reachable from m and m is directly density-reachable from p . However, p is not density reachable from q because q is not a core object.
- Similarly, r and s are density-reachable from o and o is density-reachable from r . Thus, o, r , and s are all density-connected.



DBSCAN...

“How does DBSCAN find clusters?”

- Initially, all objects in a data set D are marked as “unvisited.”
- DBSCAN randomly selects an unvisited object p , marks p as “visited,” and checks whether the ϵ -neighborhood of p contains at least MinPts objects.
- If not, p is marked as a noise point. Otherwise, a new cluster C is created for p , and all the objects in the ϵ -neighborhood of p are added to a candidate set, N .
- DBSCAN iteratively adds to C those objects in N that do not belong to any cluster. In this process, for an object p_0 in N that carries the label “unvisited,” DBSCAN marks it as “visited” and checks its ϵ -neighborhood. If the ϵ -neighborhood of p_0 has at least MinPts objects, those objects in the ϵ -neighborhood of p_0 are added to N .
- DBSCAN continues adding objects to C until C can no longer be expanded, that is, N is empty. At this time, cluster C is completed, and thus is output.
- To find the next cluster, DBSCAN randomly selects among the remaining ones. The clustering process continues until all objects are visited.



Algorithm: DBSCAN: a density-based clustering algorithm.


Input:

- D : a data set containing n objects,
- ϵ : the radius parameter, and
- $MinPts$: the neighborhood density threshold.

Output: A set of density-based clusters.

Method:

- (1) mark all objects as **unvisited**;
- (2) **do**
- (3) randomly select an unvisited object p ;
- (4) mark p as **visited**;
- (5) **if** the ϵ -neighborhood of p has at least $MinPts$ objects
- (6) create a new cluster C , and add p to C ;
- (7) let N be the set of objects in the ϵ -neighborhood of p ;
- (8) **for** each point p' in N
- (9) **if** p' is **unvisited**
- (10) mark p' as **visited**;
- (11) **if** the ϵ -neighborhood of p' has at least $MinPts$ points,
 add those points to N ;
- (12) **if** p' is not yet a member of any cluster, add p' to C ;
- (13) **end for**
- (14) output C ;
- (15) **else** mark p as **noise**;
- (16) **until** no object is **unvisited**;



DBSCAN

- If a spatial index is used, the computational complexity of DBSCAN is $O(n \log n)$, where n is the number of database objects. Otherwise, the complexity is $O(n^2)$.
- With appropriate settings of the user-defined parameters, and MinPts, the algorithm is effective in finding arbitrary-shaped clusters.

